



Age: 24

CONTACT

+33 7 67 32 20 20

Ile de France

+90 533 283 2525

kaandisli@gmail.com

<https://github.com/KaanDisli>

HARD SKILLS

- **Programming** : Python, C, Java, OCaml, Shell, SQL, .NET, JS, PyTorch.
- **Cybersecurity/Network Security** : Hackthebox, TryHackMe, Rootme, PortSwigger (<https://www.root-me.org/nathan-nut>, <https://tryhackme.com/p/kaandisli>)
- **Tools**: Burp Suite, DirBuster, SQLMap, Nmap, Wireshark, Fortinet
- **Familiar with OWASP top 10**: XSS, CSRF, directory traversal, SQL injection, file upload, JavaScript deobfuscation, RSA Encryption.

LANGUAGES

- French - Fluent
- English - Fluent
- Turkish - Fluent

INTERESTS

Economics, geopolitics, brazilian jiu-jitsu



Kaan DISLI

GRADUATE STUDENT IN ARTIFICIAL INTELLIGENCE

Education

Master's in Computer Science, Artificial Intelligence at Sorbonne University (Paris) 2024-2026

Probabilistic statistics, Markov chains, linear programming, PROLOG, PyQt, Model-based reinforcement learning, Probabilistic Turing Machines, Convolutional Neural Networks, Recurrent Neural Networks, RAG, Langchain, Petri auditing tool (Anthropic).

Bachelor's in Computer Science at Sorbonne University (Paris) 2021-2024

Artificial Intelligence for Games and Data Science - Introduction to Cryptography - Web Development Technologies: HTML/CSS, PHP, Node.js, React.js, MongoDB - Data Structures, Databases, Algorithms, Networks, and MIPS Computer Architecture.

Scientific Baccalaureate (with highest honors)

PROFESSIONAL EXPERIENCE

AI Engineer Intern 02/2026 - 08/2026 | BNP Paribas

- Designed ML pipelines for client revenue analysis, using ensemble models (CatBoost etc...) for forecasting with robust anomaly detection algorithms (IsolationForest).
- Architected the semantic matching process for data migration, building a custom NLP pipeline utilizing LLMs and BCE2 embeddings.
- Built an AI Agent pipeline from scratch using LangGraph to automatically generate PowerPoint presentations from user prompts based on business constraints and won the BNP "Changemaker" innovation challenge.

Independent Contractor / Advanced Coder AI Training 11/2024- 04/2025 | Outlier AI

- Develop and refine complex prompts to probe and challenge the reasoning capabilities of LLM's, with a focus on programming and algorithmic problem-solving. Identify model failures, provide detailed corrections, and supply high-quality feedback to improve model accuracy and robustness.

Software Engineer Intern 06/2024 - 08/2024 | Simbrella

- Back-end API development and preparation of unit and integration tests. Telegram bot for whitelisting and assigning virtual machines on vSphere ESXi for developers using pyVmomi and PostgreSQL.

PROJECTS

GPT-2 Mini – Implementation from scratch & Fine-Tuning

- Built a complete GPT-2 Small (124M params) replica in PyTorch: tokenizer, positional embeddings, multi-head self-attention, residual connections, feed-forward blocks, layer-norm, and weight-tying output head.
- Utilized AdamW, LR schedule, gradient clipping, and mixed-precision for training optimization. Tested the base model with the HellaSwag benchmark.
- Added an additional Supervised Fine-Tuning (SFT) pipeline with instruction/response format
- Built a Retrieval-Augmented Generation (RAG) pipeline for natural language question answering over PDF documents with LanceDB vector database, document chunking and response evaluator model.

LLM Fine-Tuning for Spelling Correction (Kumru-2B)

- Fine-tuned the Kumru-2B Large Language Model on a custom dataset to optimize its character-level reasoning and letter-counting capabilities.
- Implemented LoRA, reducing the trainable parameter footprint to ~33.9M to enable efficient local training.
- Achieved a 98.1% accuracy rate (up from a 16.9%) in the target task, while analyzing the trade-offs in the model's generalized capabilities.

Mechanistic Interpretability Research — Sparse Autoencoders & Jacobian Lens

- Built PyTorch interpretability pipelines for Phi-2 and Qwen models, using forward hooks to extract MLP/residual-stream activations, train L1 sparse autoencoders, and fit layerwise Jacobian lenses.
- Analyzed multi-hop reasoning by observing J-lens token trajectories; tested whether decoded concepts were causally relevant through activation steering, representation swapping, and ablation experiments.
- Fine-tuned backdoors on Qwen3.5-2B and traced trigger emergence across layers using teacher forcing, and token-rank analysis.